

Overview

- 1 Processes
- 2 Creating/Deleting Unix Processes
- 3 Inter-Process Communication
- 4 Posix Thread Programming
- 5 Java Thread Programming
- 6 Virtualization

Introduction

Distributed Programming

Distributed programming is about processes which communicate over a network

- Obvious requirement: good knowledge in how to handle processes locally

Processes

Basic idea

We build **virtual processors** in software, on top of physical processors:

- **Processor:** Provides a set of instructions along with the capability of automatically executing a series of those instructions.
- **Thread:** A minimal software processor in whose • a series of instructions can be executed. Saving a thread context implies stopping the current execution and saving all the data needed to continue the execution at a later stage.
- **Process:** A software processor in whose context one or more threads may be executed. Executing a thread, means executing a series of instructions in the context of that thread.

Threads and Distributed Systems

Multithreaded Web client

Hiding network latencies:

- Web browser scans an incoming HTML page, and finds that **more files need to be fetched**
- **Each file is fetched by a separate thread**, each doing a (blocking) HTTP request.
- As files come in, the browser displays them.

Multiple request-response calls to other machines (RPC)

Hiding network latencies:

- A client does several calls at the same time, each one by a different thread.
- It then waits until all results have been returned.
- Note: if calls are to different servers, we may have **a linear speed-up**.

Processes

- One process is made of:
 - A process identifier (an integer)
 - One executing program
 - Memory used to execute the program
 - One program counter (indicates where in the program the process currently is)
 - A number of signal handlers (tells the program what to do when receiving signals)
- One process can determine what is its own identifier:

```
#include <sys/types.h>
#include <unistd.h>
pid_t getpid(void);
```

- It can also determine the identifier of its parent:

```
pid_t getppid(void);
```


Processes

- Computers can execute programs
- A process is one instance of a program while it is executing
 - For example: I am currently using `acoread` to display these slides
 - At the same time, there are many other programs executing: they are other processes
- The same program can be executed multiple times in parallel
 - Somebody else may log on my computer and start `acoread`
 - These are several separate processes, executing the same program
- A process can only be created by another process
 - E.g., when I type a command, my shell process will create an `acoread` process

The fork() Primitive

- There is exactly one way of creating a process in Unix:

```
#include <sys/types.h>
#include <unistd.h>
pid_t fork(void);
```



The fork() Primitive

- The child process is an exact copy of the parent
 - It is running the same program
 - Its program counter is at the same position within the program
i.e., just after the fork() call
 - Its memory area is an exact copy of the parent's memory
 - Signal handlers and file descriptors are copied too
- There is one way of distinguishing between the two processes:

The `fork()` Primitive

- There is one way of distinguishing between the two processes:
 - `fork()` return 0 to the child process
 - `fork()` returns the child's pid to the parent ($pid > 0$)
 - `fork()` return -1 if an error occurred
- Most often, programs need to check the return value from `fork()`

The fork() Primitive

- After testing for *fork()* 's return value, the two processes may diverge:

```

pid_t pid;
pid = fork();
if (pid < 0)
{
    perror("Fork_error"); exit(1);
}
if (pid == 0)
{
    printf("I_am_the_child_process\n");
    while (1) putchar("c");
}
else {
    printf("I_am_the_parent_process\n");
    while (1) putchar("p");
}

```

The Fork of Death

- You must be careful when using *fork()*!
 - It is very easy to create very harmful programs
- What does the following program do?

```
while (1)
{
    fork();
}
```

The `exec()` Primitives

- With `fork()`, children always execute the same program as their parent
- `exec()` allows one process to switch to another program
 - All code/data for that process are destroyed
 - Environment variables and file descriptors are kept identical
 - New code is loaded, then started (from the beginning)

The `exec()` Primitives

- There are 4 versions of the `exec()` primitive:
 - Depending if the program is searched in `$PATH` or if it is specified by an absolute path
 - Depending if program parameters are passed as parameters to `exec()` or as a single array:

	absolute path.	search in <code>\$PATH</code>
n parameters	<code>exec()</code>	<code>execp()</code>
1 array	<code>exec()</code>	<code>execvp()</code>

- `exec()` example :

```
exec("/bin/ls", "ls", "-l", "/home/wondimagegn", NULL);
```

- `execv()` example :

```
char *params[4] = {"ls", "-l", "/home/wondimagegn", NULL};
execv("/bin/ls", params);
```

Stopping a Process

- A process stops when:
 - The `main()` function returns or
 - The program calls `exit()`
- A process cannot directly kill another process
 - It can only send a signal to it
 - Depending on the destination's signal handlers, it may stop (or do something else)
 - By default, the signal `SIGINT` stops the process. This is what happens when you type `C`

Signal Handling

- Signals can be sent by the system to a process
 - SIGSEGV: segmentation fault (non-authorized memory access)
 - SIGBUS: bus error (non-aligned memory access)
 - SIGPIPE: you tried to write in a pipe with no reader
 - SIGCHLD: one of your children process has stopped
 - There are two generic signals to be used by user programs: SIGUSR1 and SIGUSR2
 - You can get a complete list of signals with `kill -l`

Signal Handling

- To setup a signal handler:

```
#include <signal.h>
void (*signal(int signum, void (*sighandler)(int)))(int);
```

- Example:

```
#include <signal.h>
void myhandler(int sig) {
    printf("I received signal number %d!\n", sig);
    signal(sig, myhandler);
}
int main () {
    void (*oldhandler)(int);
    oldhandler = signal(SIGINT, myhandler);
    signal(SIGUSR1, myhandler);
    ...
}
```


Inter-Process Communication

- In general, processes cannot influence each other
 - Each process is executed in isolation from the others
 - For example, one process cannot write into the memory of another process. . .
- But you can decide to make several processes communicate together:
 - Send a signal to another process
 - Pipes: a communication channel to transfer data
 - Shared memory: the same memory area is accessible to multiple processes
 - Semaphores: perform synchronization (e.g., to regulate access to shared memory)
- All these IPCs only work between processes of the same computer!

Pipes

- A pipe is a unidirectional communication channel between processes
 - All data written into it can be read at the other end
 - If you need bidirectional communication, use two separate pipes!
- To create a pipe: `pipe()`

```
#include <unistd.h>
int pipe(int fd[2]);
```

Pipes

- Pipes are often used in combination with *fork()*

```

int main() {
int pid, fd[2];
char buf[64];

if (pipe(fd)<0) exit(1);
pid = fork();
if (pid==0)      /* child */
{
close(fd[0]);
/* close reader */
write(fd[1], "hello ,_world!" ,14);
}
else /* parent */
{
close(fd[1]); /* close writer */
if (read(fd[0], buf,64) > 0)
printf("Received:_%s\n", buf);
waitpid(pid, NULL, 0);
}
}

```

Pipes

- Pipes are used for example for shell commands like:

```
#include <unistd.h>

sort foo.txt | uniq | wc
```

- Pipes can only link processes which have a common ancestor
 - Because children processes inherit the file descriptors from their parent

Shared Memory

- Shared memory allows multiple processes to have direct access to the same memory area
 - They can interact through the shared memory segment
- Shared memory segments must be created, then attached to a process to be usable. Then, you must detach and destroy it afterwards.
- When using shared memory, you must be very careful about race conditions, and solve them thanks to semaphores.

Shared Memory

- To destroy a segment:

```
#include <sys/ipc.h>
#include <sys/shm.h>
int shmctl(int shmid, int cmd, struct shm_id *buf);
```

- Shared memory segments stay persistent even after all processes have died!
 - You must destroy them in your programs
 - The `ipcs` command shows existing segments (and semaphores)
 - You can destroy them by hand with: `ipcrmshm < id >`

Shared Memory

- To create a shared memory segment:

```
#include <sys/ipc.h>
#include <sys/shm.h>
int shmget(key_t key, int size, int shmflg);
```

- **key:** rendezvous point key=IPC_PRIVATE if it will be used by children processes
- **size:** size of the segment in bytes
- **shmflg:** options (access control mask)
- **Return value:** a shm identifier (or -1 for error)

Shared Memory

- To **attach** a shared memory segment:

```
#include <sys/types.h>
#include <sys/shm.h>
void *shmat(int shmid, const void *shmaddr, int shmflg);
```

- **shmid**: shared memory identifier (returned by shmget)
 - **shmaddr**: address where to attach the segment, or NULL if you don't care
 - **shmflg**: options (access control mask)
- To **detach** a shared memory segment:

```
int shmdt(const void *shmaddr);
```

- **shmaddr**: segment address
- **Attention: shmdt() does not destroy the segment!**

Shared Memory

- To destroy a segment:

```
#include <sys/ipc.h>
#include <sys/shm.h>
int shmctl(int shmid, int cmd, struct shm_id *buf);
```

- Shared memory segments stay persistent even after all processes have died!
 - You must destroy them in your programs
 - The `ipcs` command shows existing segments (and semaphores)
 - You can destroy them by hand with: `ipcrmshm < id >`

Shared Memory

Example

```
int main() {
    int shmid = shmget(IPC_PRIVATE, sizeof(int), 0600);
    int *shared_int = (int *) shmat(shmid, 0, 0);
    *shared_int = 42;
    if (fork()==0) {
        printf("The_value_is: %d\n", *shared_int);
        *shared_int = 12;
        shmdt((void *) shared_int);
    }
    else {
        sleep(1);
        printf("The_value_is: %d\n", *shared_int);
        shmdt((void *) shared_int);
        shmctl(shmid, IPC_RMID, 0);
    }
}
```

Threads vs. Processes

- Multi-process programs are expensive:
 - `fork()` needs to copy all the process' memory, etc.
 - Inter-process communication is hard
- Threads: "lightweight processes"
 - One process contains several "threads of execution"
 - All threads execute the same program (but can be at different stages within it)
 - All threads share process instructions, global memory, open files and signal handlers
 - But each thread has its own thread ID, stack, program counter and stack pointer, `errno` and signal mask
 - There are special synchronization primitives between threads of the same process

Threads in C and Java

- Threads in C:

- Posix threads (pthreads) are standard among Unix systems
- The operating system must have special support for threads
- Programs must be linked with -lpthread

- Threads in Java:

- Threads are a native feature of Java: every virtual machine has thread support
- They are portable on any Java platform
- Java threads can be mapped to operating system threads (native threads) or emulated in user space (green threads)

Creating a Pthread

- To create a pthread:

```
#include <pthread.h>
int pthread_create(pthread_t *thread, pthread_attr_t *attr,
void *(*start_routine)(void *), void *arg);
```

- thread: thread id
- attr: attributes (i.e., options)
- start routine: function that the thread will execute
- arg: parameter to be passed to the thread

- To initialize a pthread attribute:

```
int pthread_attr_init(pthread_attr_t *attr);
```

Stopping a Pthread

- A pthread stops when:
 - Its process stops
 - Its parent thread stops
 - Its start routine function returns
 - It calls pthread exit:

```
#include <pthread.h>
void pthread_exit(void *retval);
```

- Like processes, stopped threads must be waited for:

```
#include <pthread.h>
int pthread_join(pthread_t th, void **thread_return);
```

Pthread Create/Delete Example

```
#include <pthread.h>
void *func(void *param) {
    int *p = (int *) param;
    printf("New_thread: _param=%d\n", *p);
    return NULL;
}

int main() {
    pthread_t id;
    pthread_attr_t attr;
    int x = 42;
    void pthread_exit(void *retval);

    pthread_attr_init(&attr);
    pthread_create(&id, &attr, func, (void *) &x);
    pthread_join(id, NULL);
}
```

Detached Threads

- A “detached” thread:
 - Does not need to be pthread join()ed
 - Does not stop when its parent thread stops
- By default, threads are “joinable” (i.e., “attached”)
- To create a detached thread, set an attribute before creating the thread:

```
pthread_t id;  
pthread_attr_t attr;  
  
pthread_attr_init(&attr);  
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);  
pthread_create(&id, &attr, func, NULL);
```

- You can also detach a thread later with pthread detach()

Pthread Synchronization with Mutex

- Pthreads have two synchronization concepts: mutex and condition variables
- Mutex: mutual exclusion

```
#include <pthread.h>
int pthread_mutexattr_init(pthread_mutexattr_t *attr);
int pthread_mutex_init(pthread_mutex_t *mutex,
const pthread_mutexattr_t *mutexattr);
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

Pthread Synchronization with Mutex

- Pthreads have two synchronization concepts: mutex and condition variables
- Mutex: mutual exclusion

```
pthread_mutex_t array_mutex;  
int add_elem(int elem) {  
    int n;  
    pthread_mutex_lock(&array_mutex);  
    if (nbelems==32) { pthread_mutex_unlock(&array_mutex); return -1; }  
    array[nbelems++] = elem;  
    n = nbelems;  
    pthread_mutex_unlock(&array_mutex);  
    return(n);  
}  
int main() {  
    pthread_mutexattr_t attr;  
    pthread_mutexattr_init(&attr);  
    pthread_mutex_init(&array_mutex, &attr);  
    ...  
    pthread_mutex_destroy(&array_mutex);  
}
```

Creating Java Threads

- A Java thread is a class which inherits from Thread
- You must overload its run() method:

```
public class MyThread extends Thread
{
    private int argument;
    MyThread(int arg) {
        argument = arg;
    }
    public void run() {
        System.out.println("New thread started! arg=" + argument);
    }
}
```

- To start the thread:

```
MyThread t = new MyThread(42);
t.start();
```

Stopping Java Threads

- A Java thread stops when its `run()` method returns
- You do not need to `join()` for a Java thread to finish

```
MyThread t = new MyThread(42);  
t.start();  
...  
t.join();
```

Java Thread Synchronization with Monitors

- A monitor is similar to a mutex:

```
public class AnotherClass {  
    synchronized public void methodOne() { ... }  
    synchronized public void methodTwo() { ... }  
    public void methodThree()  
    { ... }  
}
```

- Each object contains one mutex, which is locked when entering a synchronized method and unlocked when leaving
 - Two synchronized methods from the same object cannot be executing simultaneously
 - Two different objects from the same class can be executing the same synchronized method simultaneously

Java Thread Synchronization with Monitors

- So, the previous class is equivalent to:

```
public class AnotherClass {  
    private Mutex mutex;  
    public void methodOne()  
    { mutex.lock(); ...; mutex.unlock(); }  
    public void methodTwo()  
    { mutex.lock(); ...; mutex.unlock(); }  
    public void methodThree() { ... }  
}
```

...except that the `Mutex` class does not exist!

Condition Variables in Java

- There is no real condition variable in Java
- But you can explicitly block a thread
 - All Java classes inherit from class Object:

```
class Object {  
    void wait();  
    /* blocks the calling thread */  
    void notify();  
    /* unblocks one thread blocked in this object */  
    void notifyAll(); /* unblocks all threads blocked in the object */  
    ...  
};
```

- wait(): causes current thread to wait until another thread invokes the notify() method or the notifyAll() method for this object.
- notify(): wakes up a single thread that is waiting on this object's monitor.
- notifyAll(): wakes up all threads that are waiting on this object's monitor.


Condition Variables in Java

- This means that each object contains one (and no more) condition variable
- The `wait()`, `notify()` and `notifyAll()` methods must be called inside a monitor

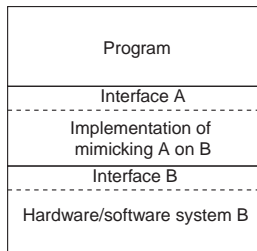
Virtualization

Virtualization is becoming increasingly important:

- Hardware changes faster than software
- Ease of portability and code migration
- **Isolation** of failing or attacked components



(a)



(b)

Architecture of VMs

Virtualization

Virtualization can take place at very different levels, strongly depending on the **interfaces** as offered by various systems components:

